

Konstantinos Voudouris

konstantinos.voudouris@dsit.gov.uk | [LinkedIn](#) | [GitHub](#) | [Scholar](#) | [Website](#)

EXPERIENCE

Research Scientist, Cognitive Science

London, UK

Alignment Team—UK AI Security Institute (PI: Geoffrey Irving)

Dec 2025–Present

Responsible for executing research projects at the intersection of cognitive science and technical alignment research, and for making targeted research grants as part of UK AISI's Alignment Project (over £30,000,000 grants made).

Post-doctoral Research Scientist, Helmholtz Munich

Munich, Germany

Institute for Human-Centered AI (PI: Eric Schulz)

Jul 2024–Nov 2025

Researching and supervising students on projects at the intersection of cognitive science and machine learning.

Post-doctoral Research Assistant/Associate, University of Cambridge

Cambridge, UK

Major Transitions in the Evolution of Cognition (PIs: M. Halina, A. Barron, C. Klein)

Nov 2023–Jul 2024

Building computational models of early eukaryotic nervous systems to better understand how they involved.

COVID-19 Infection Survey Analyst

Newport, Wales

Office for National Statistics

Jul-Oct 2021

Conducted an independent project on the causes of enduring prevalence during the second wave of the COVID-19 pandemic in the UK. Contributed to weekly governmental bulletins on prevalence and transmission.

EDUCATION

PhD (Psychology, Philosophy of Science, Artificial Intelligence)

University of Cambridge

*Passed with **no corrections** (Thesis: [From Animals to Algorithms](#)).*

Oct 2020–Mar 2024

Visiting Scholar

University of California, Berkeley

Psychology Department (Sponsor: Alison Gopnik)

Jul–Sep 2023

MPhil (Biological Sciences: Psychology)

University of Cambridge

*Passed with **minor corrections***

Oct 2019–Sep 2020

MA *Cantab.* (Theoretical and Applied Linguistics)

University of Cambridge

Starred Double First, 1st in cohort

Oct 2016–Jun 2019

SELECTED PUBLICATIONS (*FIRST AUTHOR, †SENIOR AUTHOR)

Voudouris, K.*, Akata, E., & Witte, K. (forthcoming). Judge hacking in recursive debate protocols: A call for solutions.

Voudouris, K.*, Roberts-Gaal, X.* , Buhl, M., Irving, G., & Summerfield, C. (2026). AI alignment is a human problem. *AI Security Institute*. ([Pre-Print](#))

Voudouris, K.*, Thalmann, M., Kipnis, A., Hernández-Orallo, J., & Schulz, E. (under review). Measuring What AI Systems Might Do: Towards A Measurement Science in AI. ([Pre-Print](#))

Schulze-Buschoff, L. M.* , **Voudouris, K.***, Demircan, C., & Schulz, E. (2026). Can vision language models learn intuitive physics from interaction? *43rd Intl Conference on Machine Learning (ICML) 2026*. ([Pre-Print](#))

Voudouris, K.*, Akata, E.* , Fortuin, V., & Schulz, E. (2026). Measuring In-Context Function Learning in Large Language Models. *29th Annual Conference on Artificial Intelligence and Statistics (AISTATS) 2026*. ([Pre-Print](#))

Slater, B., Mecattaf, M. G., Tešić, M., Prunty, J., **Voudouris, K.†** & Cheke, L. G.† (2025). A little less conversation, a little more action, please: Investigating the physical common-sense of LLMs in a 3D embodied environment. *PRICAI 2025*. ([Pre-Print](#))

Voudouris, K.* (2025). [Analogies And The Associative-Cognitive Distinction In Comparative Psychology](#). *Biology & Philosophy*, 40, 19. ([Pre-Print](#))

Binz, M., ..., **Voudouris, K.**, ..., **Schulz, E.** (2025) [A foundation model to predict and capture human cognition](#). *Nature* 644, 1002-1009.

Voudouris, K.*, Schulze-Buschoff, L. M.* , Akata, E., Bethge, M., Tenenbaum, J., & Schulz, E. (2025) Testing the limits of fine-tuning for improving visual cognition in vision language models. *42nd Intl Conference on Machine Learning (ICML) 2025*. ([Pre-Print](#))

Pacchiardi, L., **Voudouris, K.**, Slater, B., Martinez-Plumed, F., Zhou, L., Hernández-Orallo, J., & Schellaert, W. (2025) PredictaBoard: Benchmarking LLM Score Predictability. *Findings of the ACL*. ([Pre-Print](#))

Voudouris, K.*, Cheke, L. G., & Schulz, E. (2025) [Bringing comparative cognition approaches to AI systems](#). *Nature Reviews Psychology*.

Voudouris, K.*, ..., & Cheke, L. G. (2025) [The Animal-AI Environment: A Virtual Laboratory For Comparative Cognition and Artificial Intelligence Research](#). *Behavior Research Methods*, 57:107.

Kipnis, A., **Voudouris, K.**, Buschoff Schulze, L., & Schulz, E. (2025) metabench—A sparse benchmark for reasoning and knowledge in large language models. *13th Intl Conference on Learning Representations (ICLR) 2025*. ([Pre-Print](#)).

Voudouris, K.*, Farrar, B. G., Cheke, L. G., & Halina, M. (2025) Morgan’s Canon and the Associative-Cognitive Distinction Today: A Survey of Practitioners. *Journal of Comparative Psychology* ([Pre-Registration](#); [Pre-Print](#)) ([JCP Best Paper Award Finalist](#)).

Burden, J.* , **Voudouris, K.*** , Burnell, R., Rutar, D., Cheke, L., & Hernández-Orallo, J. (2025) Inferring Capabilities From Task Performance with Bayesian Triangulation. ([Pre-Print](#)).

Voudouris, K.* , Liu, J. D., Siwinska, N., Schellaert, W., & Cheke, L. G. (2024). [Investigating Object Permanence in Deep Reinforcement Learning Agents](#). In *Proceedings of the Annual Meeting of the Cognitive Science Society (Vol. 46)*.

Voudouris, K.* , Cheke, L. G., & Halina, M. (2024). [The Future is Computational Comparative Cognition](#). *Comparative Cognition & Behavior Reviews*, 19.

Chan, A., ..., **Voudouris, K.**, ... & Maharaj, T. (2023). [Harms from Increasingly Agentic Algorithmic Systems](#). *ACM Conference on Fairness, Accountability, and Transparency 2023, Chicago, USA*. ([Pre-Print](#)).

Voudouris, K.* , Donnelly, N., Rutar, D., Burnell, R., Burden, J., Hernández-Orallo, J., & Cheke, L. G. (2022). [Evaluating Object Permanence in Embodied Agents Using the Animal-AI Environment](#). *EBeM’22: Workshop on AI Evaluation Beyond Metrics, July 25, 2022, Vienna, Austria*.

Voudouris, K.* , Crosby, M., Beyret, B., Hernández-Orallo, J., Shanahan, M., Halina, M., & Cheke, L. (2022). [Direct Human-AI Comparison in the Animal-AI Environment](#). *Frontiers in Psychology* 13:711821.

RESEARCH PROJECT SUPERVISION

AI Safety Supervisor <i>Pivotal Fellowship - 1 fellow</i>	Jun 2026–present <i>London, England</i>
AI Safety Co-Supervisor <i>LASR Labs - 4 fellows</i>	Feb 2026–May 2026 <i>London, England</i>
AI Safety Research Sponsor <i>AISI Alignment Project & Challenge Fund - 4 grants</i>	Dec 2025–present <i>London, England</i>
PhD Student Assistant Supervisor <i>Helmholtz Munich - 2 students</i>	Jul 2024–Present <i>Munich, Germany</i>
Psychology & Behavioural Sciences Dissertation Supervisor <i>Department of Psychology (Statistics, AI Evaluation, Cognitive Modelling) - 4 Students</i>	Oct 2021—Jun 2023 <i>Cambridge, UK</i>

WORKSHOPS, TUTORIALS, & SYMPOSIA

Co-organiser and co-presenter of [The Cognitive Science of AI Alignment](#) Workshop at CogSci’26 (Rio de Janeiro, Jul 2026).

Co-organiser and co-presenter of the [Measurement Layouts for Capability-oriented AI Evaluation](#) Lab at AAAI-24 (Vancouver, Feb 2024).

Co-organiser of the [Predictable AI](#) Conference in Valencia, Spain (8th March 2023).

TEACHING

Cambridge University Undergraduate Supervision: **Cognitive Neuroscience & Experimental Psychology** (Part IB, second-year; 2022-2023), **Philosophy of Cognitive Science** (Part II, third-year; 2022-2023), **Comparative Cognition** (Part II, third-year; 2022-2023), **Phonological Theory** (Part II, second- and third-year; 2020-2022). **Introduction to Linguistics** (2020, 2021; [Syllabus](#)).

TECHNICAL SKILLS

Programming Languages: Python, Julia, R, PowerShell, SQL, L^AT_EX.

Libraries and Frameworks: git, Docker, Lux.jl, MySQL, SLURM, tidyverse, torch, transformers, trl, unsloth.

Natural Languages: English (native), Greek (C1), French (B2), German (B1), Turkish (A2), Mandarin Chinese (A2).